

Deep Learning

CSC-Elective

Instructor : Dr. Muhammad Ismail Mangrio

Slides credit Dr. M Asif Khan

ismail@iba-suk.edu.pk

Week 4-5

Contents

- Loss functions
- Loss functions for regression
- Loss functions for classification
- Optimizers
- Epoch and iteration

Loss functions

- The loss function is an expression used to measure **how close the predicted value is to the actual value.**
- This expression outputs a value called loss, which tells us the performance of our model.
- By **reducing this loss value** in further **training**, the model can be **optimized** to output values that are **closer to the actual values**.
- There are **two classes** of Loss functions based on the type of learning task
 - **Loss functions for Regression models:** predict continuous values.
 - **Loss functions for Classification models:** predict the output from a set of finite categorical values.

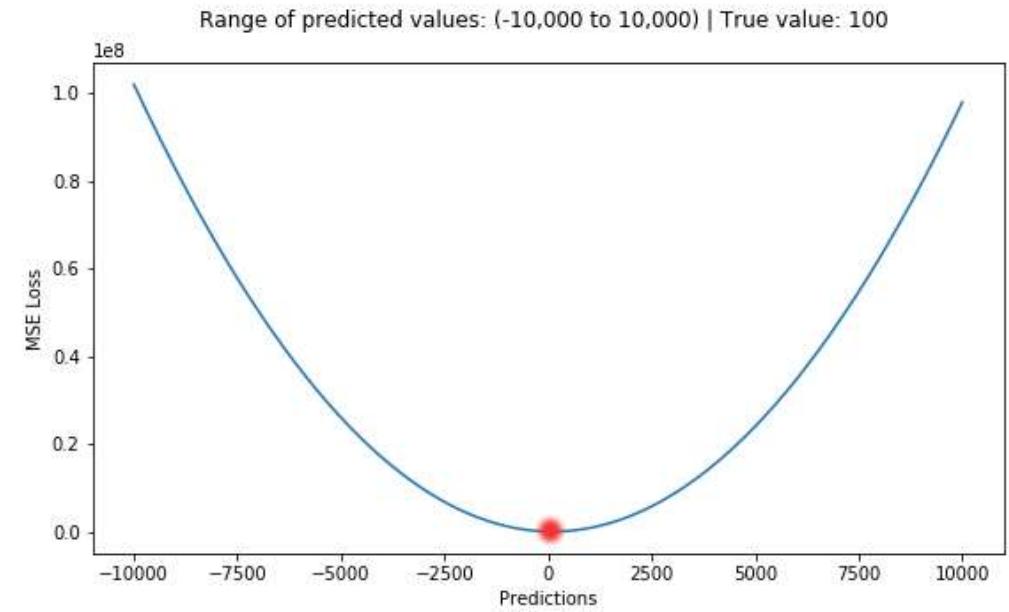
Loss functions (for Regression models)

- In Regression we predict a continuous value.
- Following are the common loss functions for Regression models:
 - **MSE** (Mean Squared Error)
 - **MAE** (Mean Absolute Error)
 - **Huber Loss function**

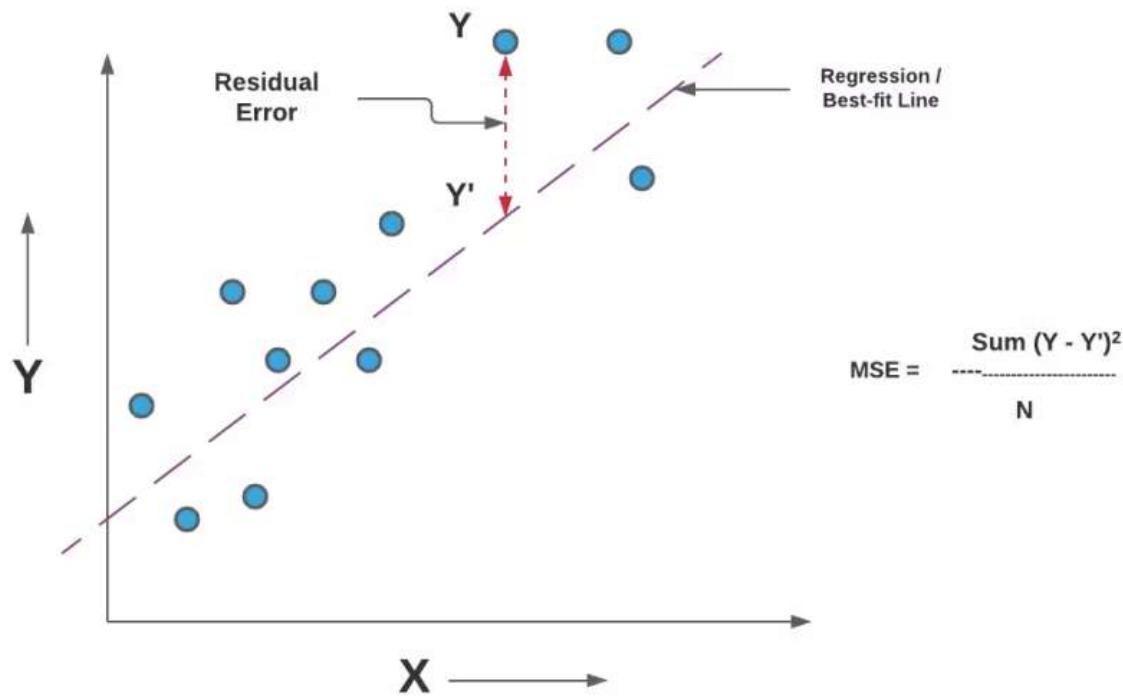
Loss functions (for Regression models, MSE)

- It is the **Mean of Square of Residuals** for all the data points in the dataset.
- **Residuals** is the difference b/w the actual and the predicted prediction by the model.
- **Squaring of residuals** is done to convert negative values to positive values.
- **Squaring also gives more weightage to larger errors.**
- When the cost function is far away from its minimal value, squaring the error will **penalize the model**
- MSE is **sensitive to outliers**.
- It is also called Quaratic Loss or L2 Loss.
- Good for gradient descent and differentiable

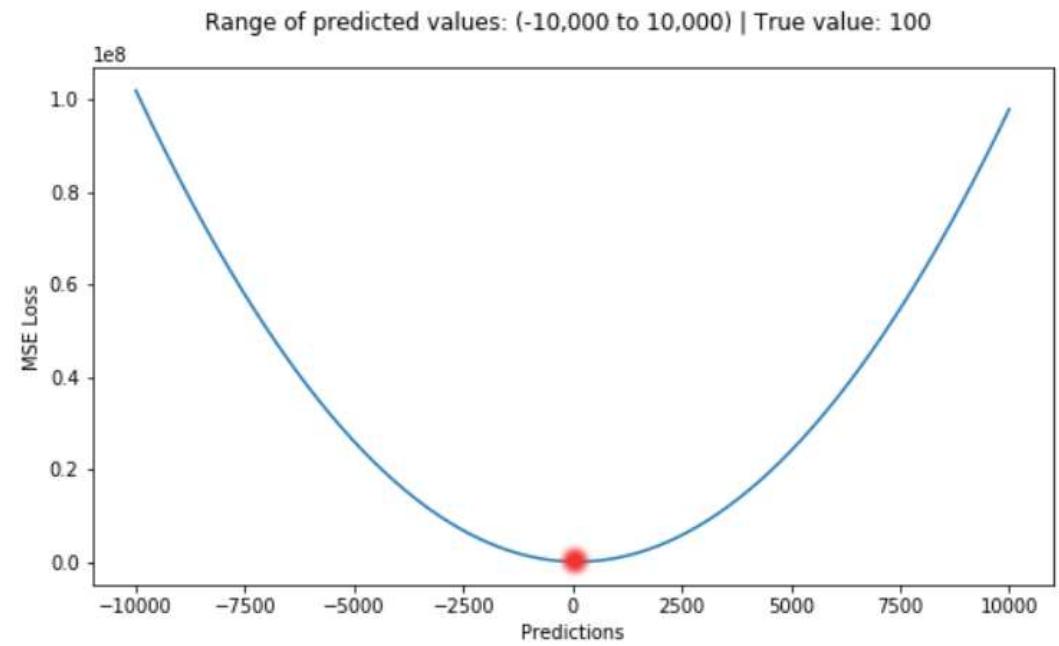
$$MSE = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$



Loss functions (for Regression models, MSE)



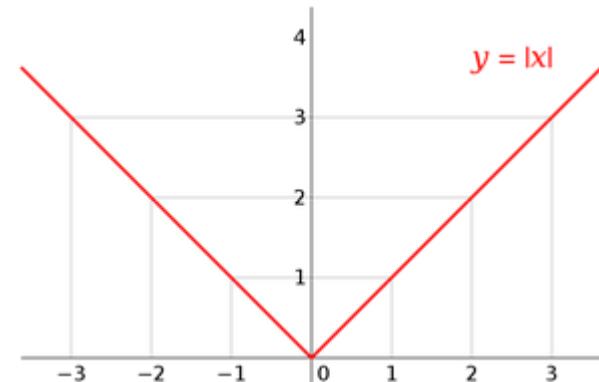
$$MSE = \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{n}$$



Loss functions (for Regression models, MAE)

- It is the **Mean of Absolute of Residuals** for all the datapoints in the dataset.
- The **absolute of residuals** is done to convert negative values to positive values.
- Mean is taken to make the loss function independent of number of datapoints in the training set.
- One advantage of MAE is that is **robust to outliers**, as it is **not squaring error**.
- MAE is generally less preferred over MSE as it is **harder to calculate the derivative of the absolute function** because absolute function is **not differentiable** at the minima.
- It is also called L1 loss.

$$MAE = \frac{\sum_{i=1}^n |y_i - \hat{y}_i|}{n}$$



Loss functions (for Regression models, Huber)

- It is the **combination of MSE and MAE**.
- It **takes the good properties of both** loss functions by being **less sensitive to outliers and differentiable at minima**.
- When the **error is smaller**, the **MSE part** of the Huber is utilized and when the **error is large**, the **MAE part** of Huber loss is used.
- A new hyper-parameter ' **δ** ' is introduced which **tells** the loss function **where to switch** from MSE to MAE.
- **Additional ' δ ' terms** are introduced in the loss function **to smoothen** the transition from MSE to MAE.

$$\text{Loss} = \begin{cases} \frac{1}{2} * (x - y)^2 & \text{if } (|x - y| \leq \delta) \\ \delta * |x - y| - \frac{1}{2} * \delta^2 & \text{otherwise} \end{cases}$$

Loss functions (for Classification models)

- The **Cross Entropy Loss** function is a commonly used loss function for classification problems, both in **binary and multi-class** settings.
- Measures the **difference b/w two probability distributions**: the predicted probability distribution and the true probability distribution (or one-hot encoded labels).
- **Binary Cross Entropy (BCE)** loss function
 - Used for binary classification problems
- **Categorical Cross Entry (CCE)** loss function
 - Used for multi classification problems, used with one-hot encoded vectors
- **Sparse Categorical Cross-Entropy (SCCE)** loss function:
 - Similar to categorical cross-entropy but used when the class labels are integers instead of one-hot encoded vectors.

Loss functions (for Classification models, BCE)

- Used for **binary classification tasks**, where each instance is classified into one of two classes (e.g., 0 or 1).
- For a single instance, if the true label is y and the predicted probability of the positive class is p , the **Binary Cross Entropy Loss** is:

$$\text{BCE} = -[y \log(p) + (1 - y) \log(1 - p)]$$

- y : True label (0 or 1).
- p : Predicted probability of the positive class ($0 \leq p \leq 1$).
- **Example:** Suppose you have: True label (y) = 1
- Predicted probability (p) = 0.8
- $\text{BCE} = -[1 \log(0.8) + (1-1) \log(1-0.8)] = -\log(0.8) \approx 0.223$

Loss functions (for Classification models, CCE)

- Used for **multi-class classification** tasks, where each instance can be classified into one of multiple classes.
- **Formula:** For a single instance, if the true label is a **one-hot encoded vector y** (with a 1 in the position of the true class and 0 elsewhere), and the predicted probability distribution is p , the Categorical Cross Entropy Loss is:

$$CCE = -\sum_i y_i \log(p_i)$$

- y : True label as a one-hot encoded vector.
- p : Predicted probability distribution.

Loss functions (for Classification models, CCE)

$$CCE = -\sum_i y_i \log(pi)$$

- y : True label as a one-hot encoded vector.
- p : Predicted probability distribution.
- **Example:** Suppose you have a classification problem with 3 classes, and for a given instance:
 - True label is class 2 (**one-hot encoded as [0, 1, 0]**).
 - Predicted probabilities are [0.1, 0.7, 0.2].
 - $CCE = -[0\log(0.1)+1\log(0.7)+0\log(0.2)] = -\log(0.7) \approx 0.357$

Loss functions (for Classification models, SCCE)

- **Purpose:** Used for multi-class classification tasks.
 - Application: Suitable when **target labels are integers** (**not one-hot encoded**) (e.g., 0, 1, 2 for 3 classes).
 - Model Output:
 - Produces a probability distribution over all classes.
 - Example: [0.7 (Cat), 0.2 (Dog), 0.1 (Rabbit)]
 - True Label:
 - Represented as an integer.
 - For example, 0 for Cat.
 - Loss Calculation:
 - Formula: **$-\log(p_{true})$**
 - Where p_{true} is the predicted probability for the true class.
- Example: True label: 0 (Cat)
Predicted probabilities: [0.7, 0.2, 0.1]
 $\text{Loss} = -\log(0.7) \approx 0.357$

Class Performance test

- Answer (type of problem, loss functions name optimizer) for following ANN:
- Input -> ReLU -> softmax()
- Input -> PReLU -> softmax()
- Input -> ReLU -> sigmoid()
- Input -> ReLU -> linear()
- You have 5 minutes to answer

Optimizers

- Optimizers in deep learning are algorithms **used to adjust the parameters (weights and biases)** of a neural network in order **to minimize the loss function** during training.
- Play a crucial role in the training process by determining how the model's parameters are **updated in response to the gradients computed** during back propagation.
- Most Common optimizers are:
 - Gradient Descent
 - SGD (Stochastic Gradient Descent)
 - Mini-Batch Gradient Descent
 - SGD Momentum
 - Adagrad
 - RMSprop
 - Adam (Adaptive Moment Estimation)

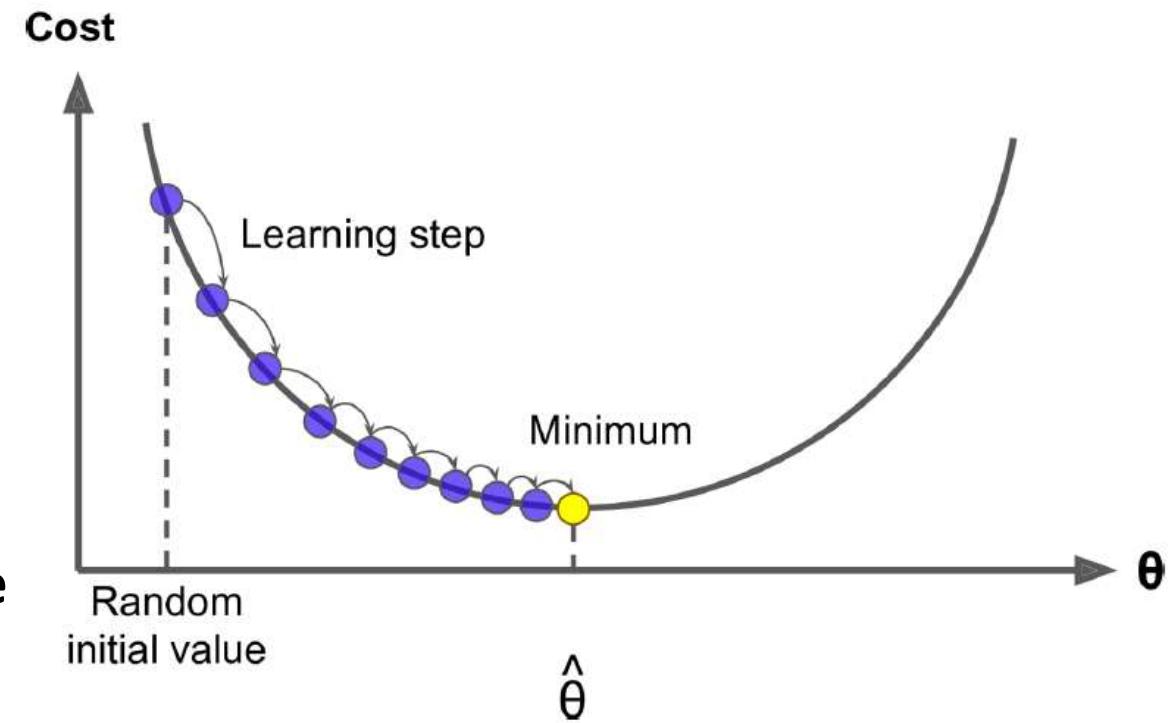
Optimizers (GD-Gradient descent revisited)

- The general idea of Gradient Descent is to **tweak parameters iteratively** in order to **minimize a cost function**.
- Suppose you are **lost in the mountains in a dense fog**; you can only **feel the slope** of the ground below your feet.
- A good strategy to get to the bottom of the valley quickly is to **go downhill in the direction of the steepest slope**.
Exactly what Gradient Descent does:



Optimizers (Gradient descent revisited)

- It measures local gradient of the error function with respect to the parameter vector θ , and goes in the direction of descending gradient.
- Once the gradient is zero, you have reached a minimum!
- Concretely, you start by filling θ with random values (this is called random initialization), and then you improve it gradually, taking one baby step at a time, each step attempting to decrease the cost function (e.g., the MSE), until the algorithm converges to a minimum (see Figure 4-3, see book).



Optimizers (Gradient descent revisited)

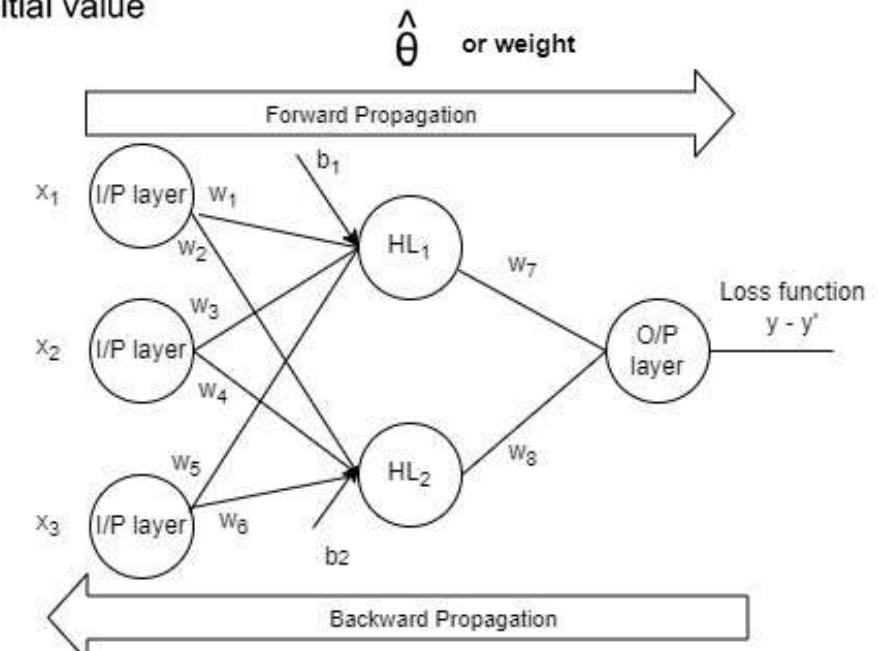
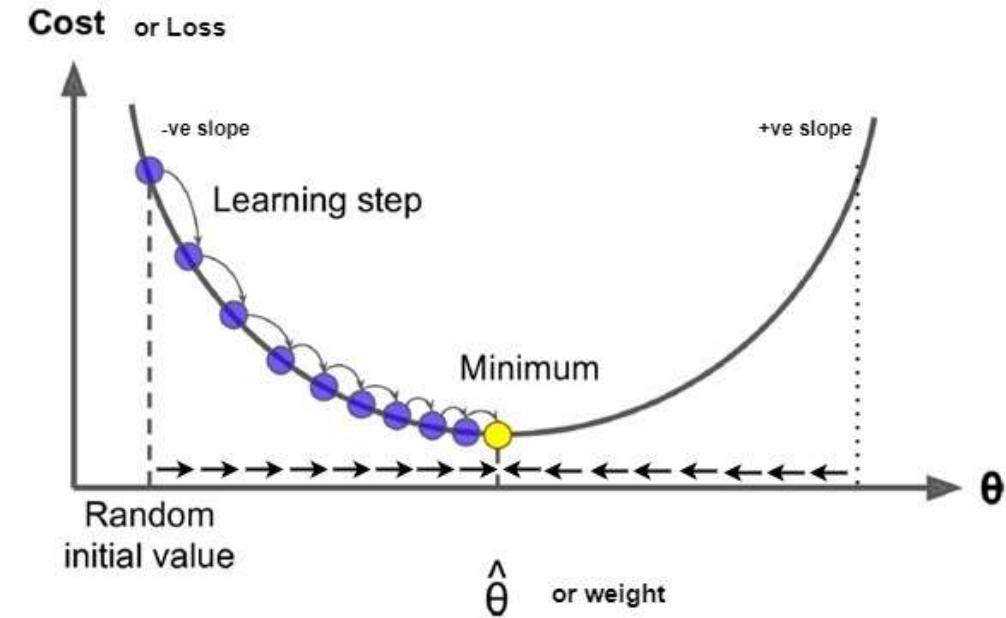
- **Weight update formula**

$$W_{new} = W_{old} - \eta \frac{\partial L}{\partial W_{old}}$$

- Here η is **learning rate**
- The **partial derivate of loss w.r.t to W_{old}** is the slope of tangent line towards minima
- For W_8 to update its weight the formula would be:

$$W_{8new} = W_{8old} - \eta \frac{\partial L}{\partial W_{8old}}$$

- The learning rate η is **kept small to avoid big jumps**, it is usually **kept 0.001**.



Optimizers (Gradient descent revisited)

- Weight update formula

$$W_{new} = W_{old} - \eta \frac{\partial L}{\partial W_{old}}$$

- From **left side** of the curve the gradient/slope is **-ve**. Therefore:

$$W_{new} = W_{old} - \eta \{-ve \text{ slope}\}$$

$$W_{new} \gg W_{old}$$

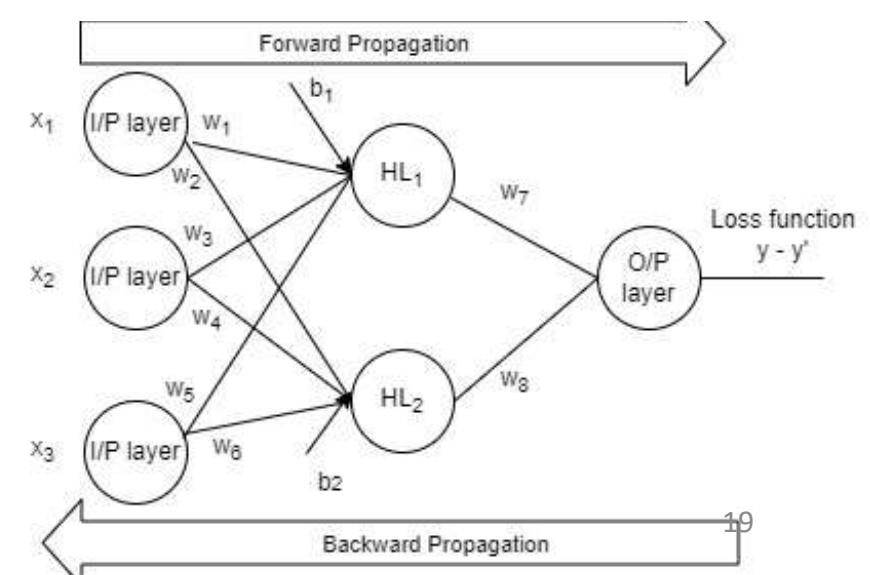
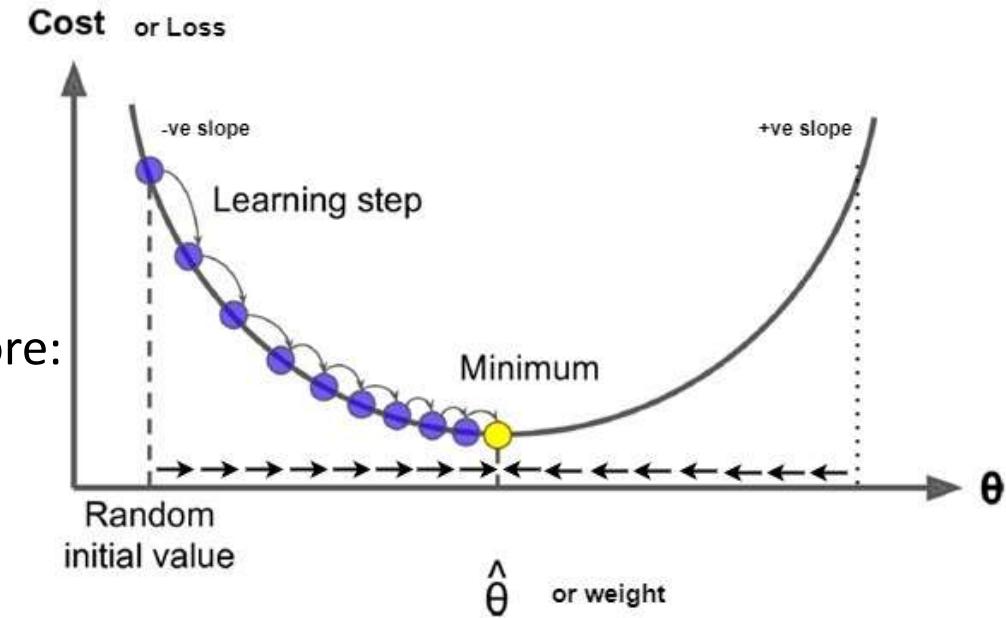
- On-contrary gradient/slope from the **right side** curve is **+ve** and:

$$W_{new} = W_{old} - \eta \{+ve \text{ slope}\}$$

$$W_{new} \ll W_{old}$$

- When updated **W reaches global minima**:

$$W_{new} = W_{old}$$



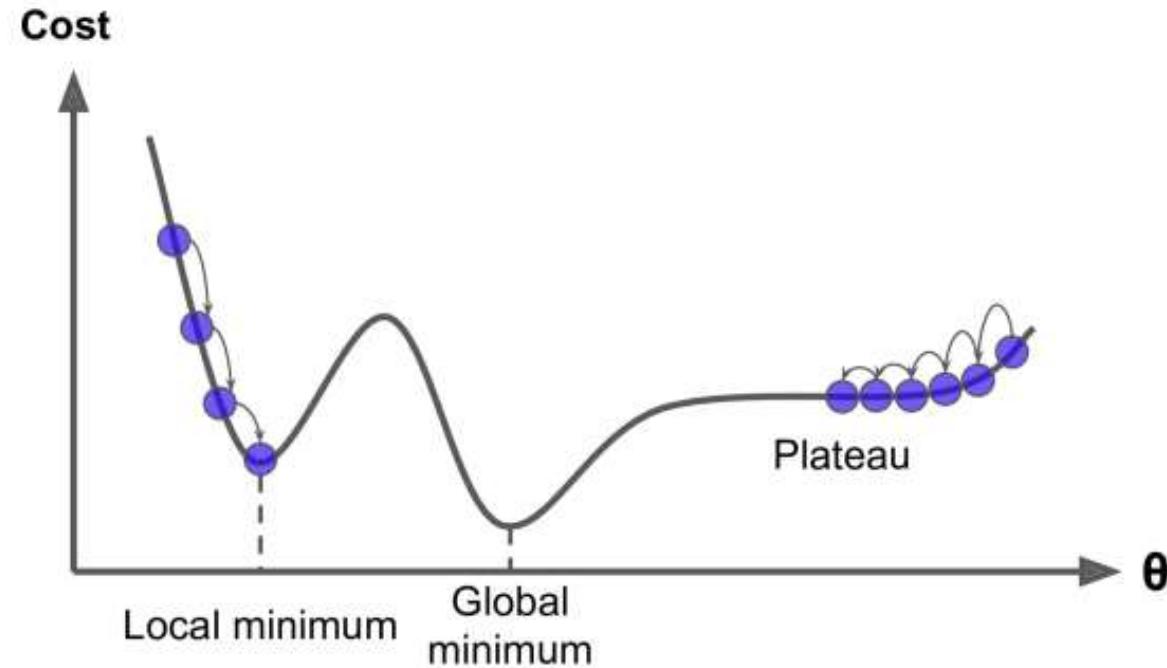
Optimizers (Gradient descent revisited)

- **Advantage**

- It will converge

- **Disadvantage**

- Needs loads of resources (e.g., RAM, CPU power etc).
- As it **needs to pass entire dataset** a once and then calculate average loss function and average gradient,
- then back propagation is applied to adjust weights. This process take **huge time**.
- It can **get stucked in local minima**



Epoch

- An epoch in ANN refers to one **complete pass through the entire training dataset**.
- It is **one full iteration over the entire training dataset**.
- **Iteration:** A single update of the model's weights, based on a batch of training data.
- **Batch:** A subset of the training data.
- Each epoch represents an opportunity for the model to **learn from the entire dataset and update its parameters** based on the computed gradients.
- Here's an example:
 - An **epoch in GD means** the model has processed every sample in the training dataset once, and the parameters have been updated based on the aggregate gradients computed from the entire dataset.
 - **Training typically involves multiple epochs** to allow the model to learn better from the data

Epoch

- **Example:** Imagine you have a dataset with 100 samples, and you are training a model with a batch size of 10.
- **Training Process for One Epoch:**
 - **Initialization:**
 - Dataset: 100 samples.
 - Batch Size: 10 samples.
 - **Epoch 1:**
 - Iteration 1: Process samples 1 to 10.
 - Iteration 2: Process samples 11 to 20.
 - Iteration 3: Process samples 21 to 30.
 - Iteration 4: Process samples 31 to 40.
 - Iteration 5: Process samples 41 to 50.
 - Iteration 6: Process samples 51 to 60.
 - Iteration 7: Process samples 61 to 70.
 - Iteration 8: Process samples 71 to 80.
 - Iteration 9: Process samples 81 to 90.
 - Iteration 10: Process samples 91 to 100.

After Iteration 10,
all 100 samples have been processed once,
completing Epoch 1.

Epoch (How many Epochs?)

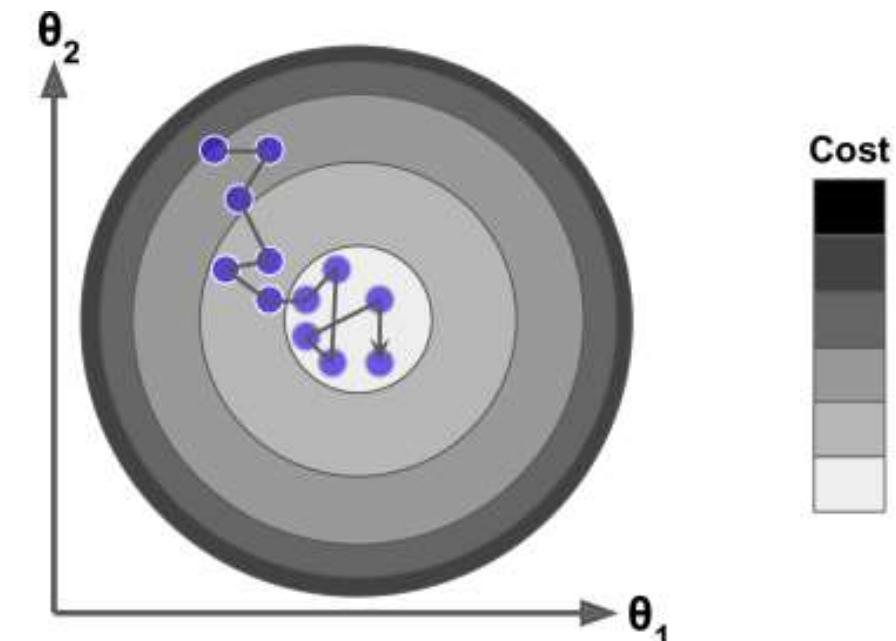
- Why **Multiple Epochs?**
- **Learning:** The model **might not learn enough** from a **single pass**.
- Multiple epochs **help the model learn patterns and reduce error**.
- **Overfitting Risk:** Too many epochs can **lead to overfitting**.
- It's essential to **monitor validation** performance.

Optimizers (SGD)

- It's a **variation of the Gradient Descent algorithm**.
- In Gradient Descent, we analyze the entire dataset in each step, which may not be efficient when dealing with very large datasets.
- To address **issue in GD**, we have **Stochastic Gradient Descent (SGD)**.
- In SGD, we **process just one example at a time** to perform a single step.
- So, if the dataset contains 10000 rows, SGD will **update the model parameters 10000 times** in a single cycle through the dataset, as opposed to just once in the case of Gradient Descent.

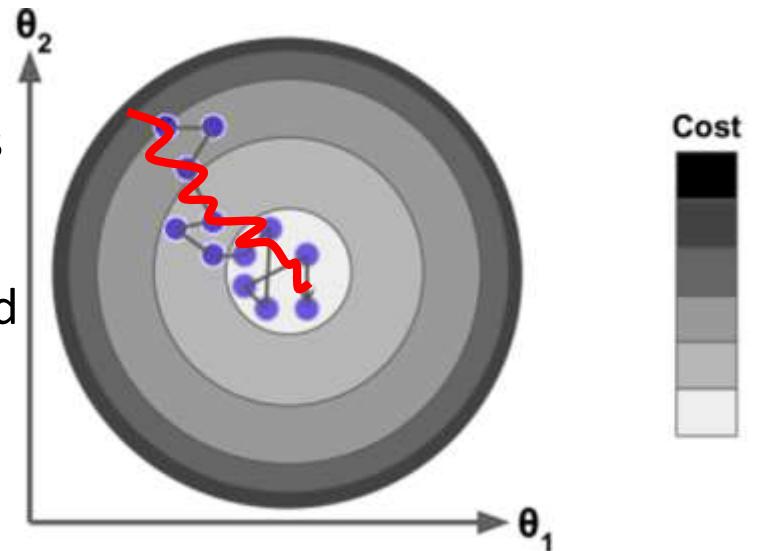
Optimizers (SGD)

- We don't need large RAM for SGD
- Here the convergence will be slower as if we had a dataset with 1 million records then weights will be updated 1 million times.
- As result it takes more time.
- Also, the gradient will be noisy.



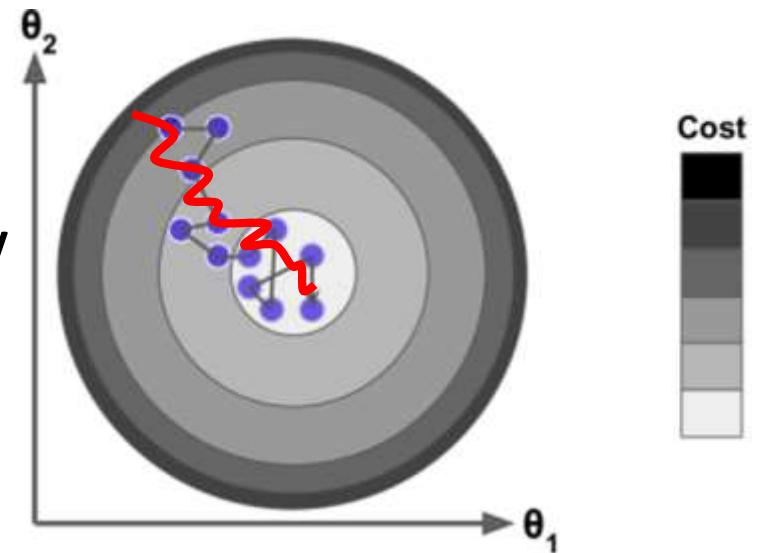
Optimizers (mini-batch SGD)

- Mini-batch stochastic gradient descent, consists of a **predetermined number of training examples**, smaller than the full dataset.
- This approach **combines the advantages GD and SGD**.
- **In one epoch**, following the creation of fixed-size mini-batches with following steps:
 - Select a mini-batch.
 - Compute the mean gradient of the mini-batch.
 - Apply the mean gradient obtained in step 2 to update the model's weights.
 - Repeat steps 1 to 2 for all the mini-batches that have been created

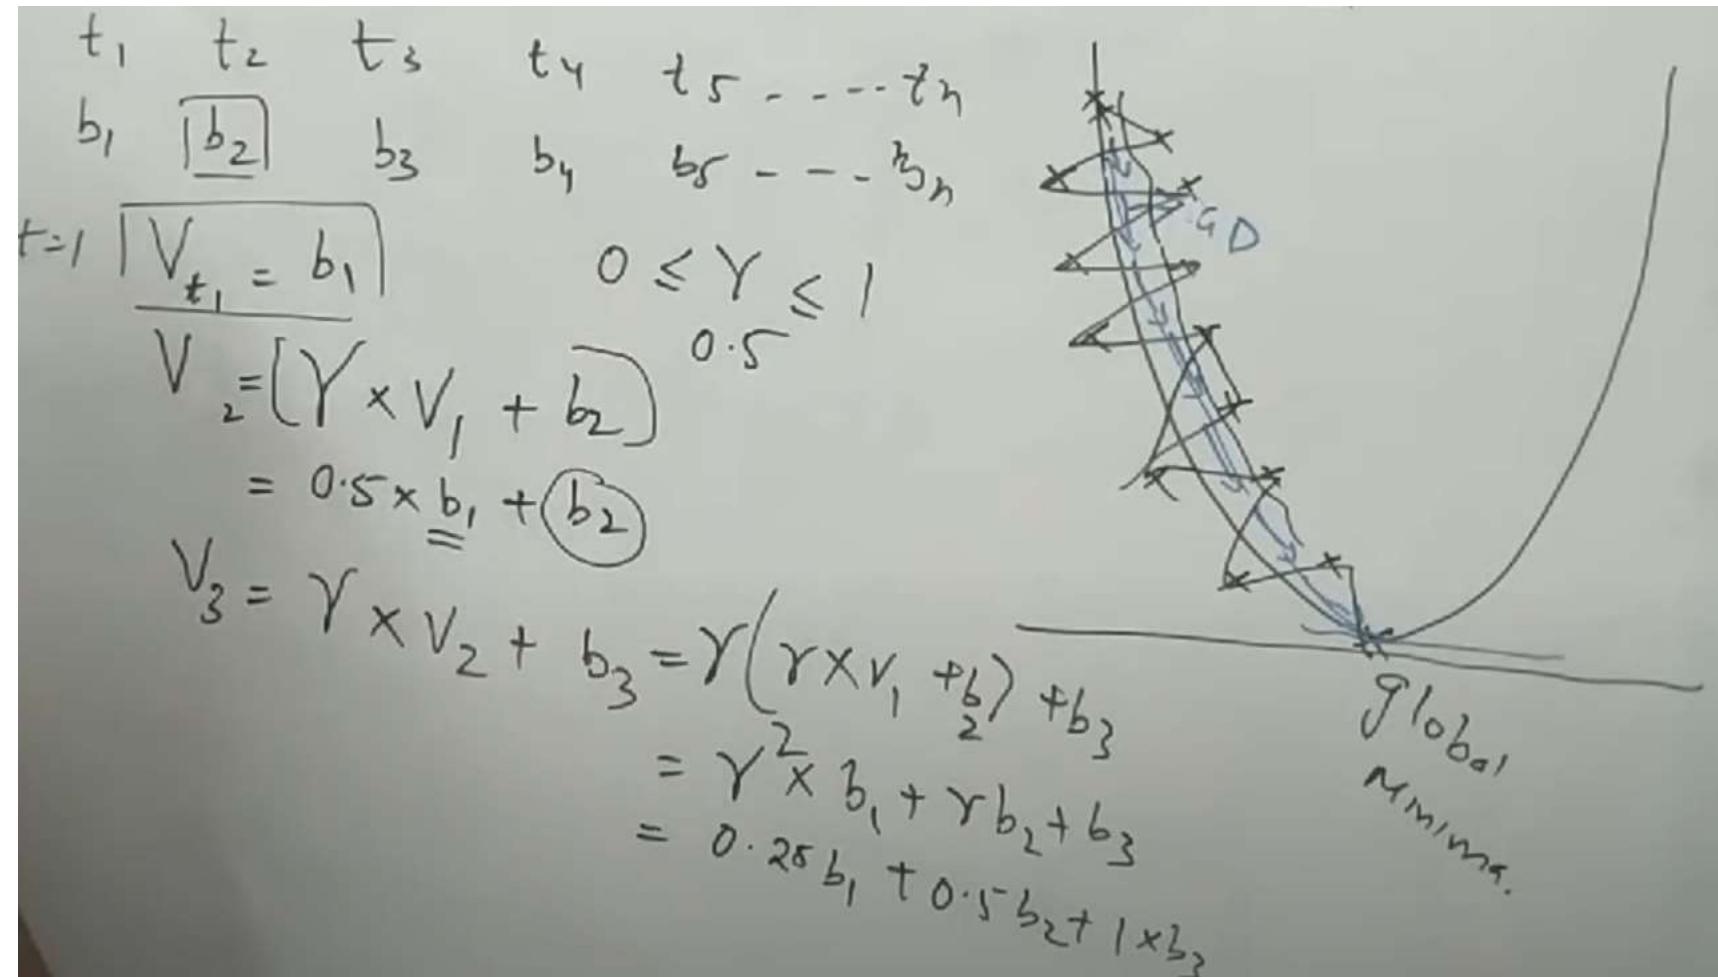


Optimizers (SGD with Momentum)

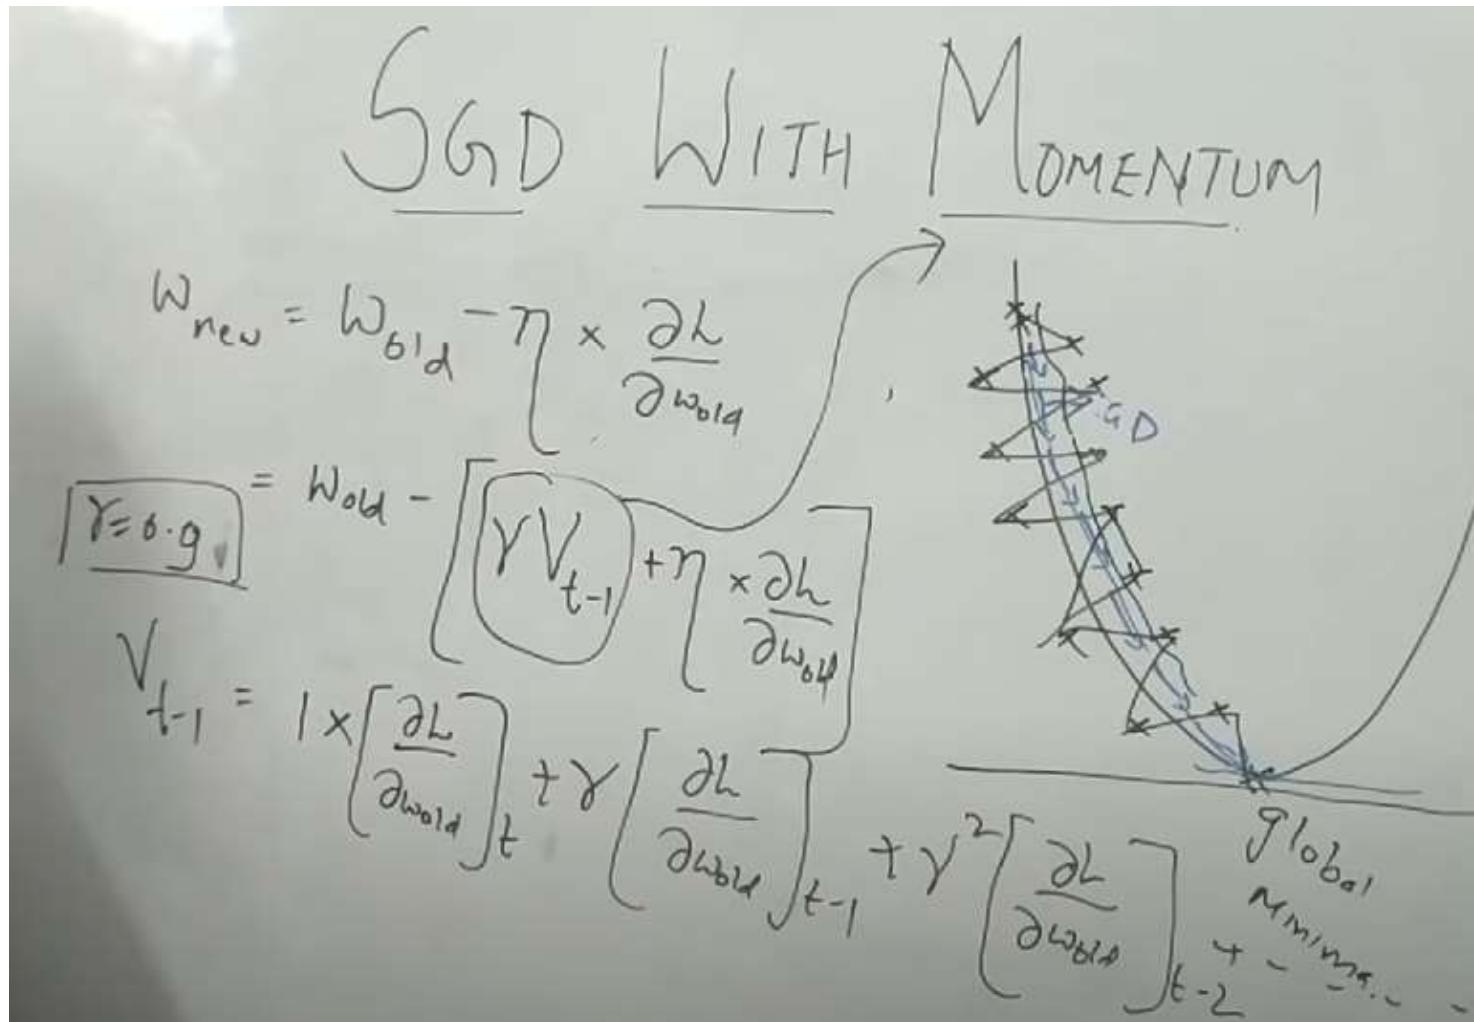
- In Stochastic Gradient Descent, we don't calculate the precise derivative of our loss function.
- Instead, we estimate it using a small batch. This results in “noisy” derivatives, which implies that we don't always move in the optimal direction.
- To address this issue, **Momentum was introduced to mitigate the noise in SGD.**
- It speeds up convergence towards the relevant direction and diminishes fluctuations in irrelevant directions.
- The concept behind Momentum involves **denoising the derivatives by employing an exponential weighting average** by assigning more weight to recent updates compared to previous ones



Optimizers (SGD with Momentum)



Optimizers (SGD with Momentum)



Optimizers (SGD with Momentum)

- Momentum works by **maintaining a velocity vector** that is updated with a combination of the previous velocity and the current gradient. The velocity vector is used to update the model parameters.

Exponential Weighted Average $v_t = \beta \times v_{t-1} + (1 - \beta) \times \nabla L(\theta)$

$$W_{new} = W_{old} - \eta v_{dw}$$

$$W_{new} = W_{old} - \eta \frac{\nabla L}{\partial W_{old}}$$

- v_t (v_{dw}) is the velocity at time t .
- β is the momentum coefficient (typically set between 0.8 and 0.99).
- v_{t-1} is the previous velocity.
- $\nabla L(\theta)$ is the gradient of the loss function with respect to the model parameters.
- We will apply v_t instead derivative of L w.r.t to old weight to update to weights

Optimizers (Adagrad)

- **Dense vs Sparse Features**

Imagine we are storing **user features**:

- **Dense features:**

[25.0, 170.5, 68.2] (Age, Height, Weight)

- **Sparse features:**

[0, 0, 1, 0, 0, 0, 0] (One-hot encoded country → Pakistan)

- **Dense features:** continuous values, mostly non-zero, used directly.
- **Sparse features:** categorical or text data, mostly zeros, stored efficiently to save space.

Optimizers (Adagrad)

- One Hot Encoding

- A method to represent **categorical data (labels, categories, classes)** as numbers so that a computer can understand them.
- Instead of giving a category a single number (like Pakistan = 1, USA = 2, UK = 3), which could confuse the model (it may think $3 > 1$),
- We use a **binary vector where only one position is "hot" (1)**, and all others are 0.

Country	One-hot Vector
Pakistan	[1, 0, 0]
USA	[0, 1, 0]
UK	[0, 0, 1]

Optimizers (Adagrad)

- One Hot Encoding

```
import torch
import torch.nn.functional as F

# Suppose we have 3 categories: 0=Pakistan, 1=USA, 2=UK
categories = torch.tensor([0, 1, 2])

# Apply one-hot encoding
one_hot = F.one_hot(categories, num_classes=3)

print(one_hot)

tensor([[1, 0, 0],
        [0, 1, 0],
        [0, 0, 1]])
```

Optimizers (Adagrad)

- Imagine you are learning to walk on different terrains:
 - On a flat road (easy), you quickly learn → so you take smaller steps later.
 - On a rocky road (hard), you need to keep taking bigger steps to adjust.
- AdaGrad does exactly this for parameters!
- Each parameter gets its own adaptive learning rate, based on how often it has been updated.

Optimizers (Adagrad)

- AdaGrad, short for adaptive gradient, signifies that the learning rates are adjusted or adapted over time based on previous gradients.
- A limitation of the previously discussed optimizers (SGD, Minibatch SGD) is the use of a **fixed learning rate** for all parameters throughout each cycle.
- This can **hinder the training features** which often exhibit **small average gradients** causing them to **train at a slower pace**.
- A **solution is to set different learning rates for each feature**, this can become complex .
- AdaGrad addresses this issue by implementing the concept that **the more a feature has been updated in the past, the less it will be updated in the future**.
- This provides an opportunity for other features, such as **sparse** features, to catch up.
- AdaGrad, as an optimizer, **dynamically adjusts the learning rate** for each parameter at every time step ‘t’.

Optimizers (Adagrad)

- ϵ is a small positive value number to avoid divide by zero error if in case $\alpha(t)$ becomes 0
- g_i is derivative of loss w.r.t weight and g_i^2 will always be +ve since its a square term, this implies that $\alpha(t) \geq \alpha(t-1)$
- One main **disadvantage** of Adagrad optimizer is that **alpha(t)** can become large as the number of iterations will increase and due to this η_t ' will decrease at the larger rate.
- This will make the old weight almost equal to the new weight which may lead to slow convergence.

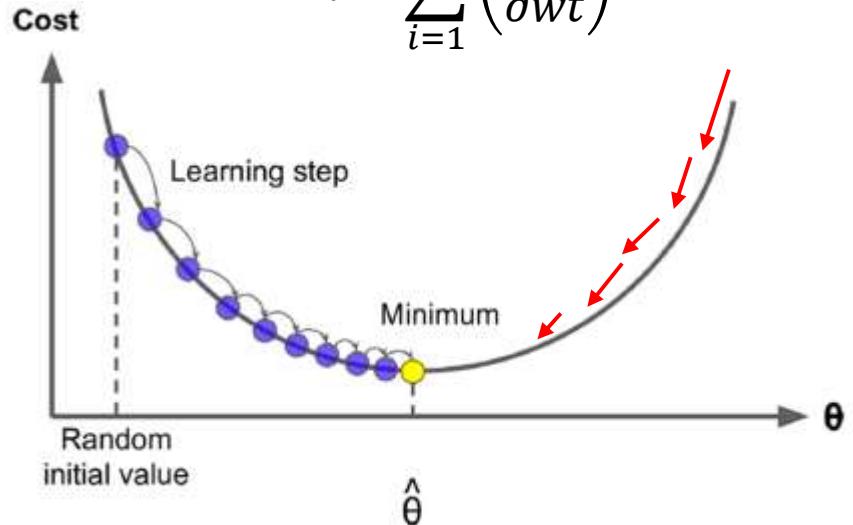
$$(w)_{\text{new}} = (w)_{\text{old}} - \eta \frac{\partial L}{\partial w(\text{old})}$$

$$w_t = w_{t-1} - \eta \frac{\partial L}{\partial w(t-1)}$$

$$w_t = w_{t-1} - \eta'_t \frac{\partial L}{\partial w(t-1)}$$

$$\eta'_t = \frac{\eta}{\sqrt{\alpha_t + \epsilon}}$$

$$\alpha_t = \sum_{i=1}^t \left(\frac{\partial L}{\partial w_t} \right)^2$$



Optimizers (RMSProp)

- RMSprop stands for **Root Mean Square Propagation**.
- RMSprop optimizer **doesn't let gradients accumulate for momentum** instead only accumulates gradients in a particular fixed window (**avoid alpha to reach big value**).
- It can be considered as an **updated version of AdaGrad** with few improvements.

$$w_t = w_{t-1} - \eta'_t \frac{\partial L}{\partial w(t-1)} \quad \eta' = \frac{\eta}{\sqrt{Sdw + \epsilon}}$$

$$Sdw_t = \beta \times Sdw_{t-1} + (1 - \beta) \times \left(\frac{\partial L}{\partial w_t} \right)^2$$

- Sdw value at initial time t will be 0
- The drawback of AdaGrad is learning rate some times can become very slow due to high α_t value.
- Here α_t is replaced by Sdw that slows the derivative of L w.r.t weight by multiplying it with $(1 - \beta)$.
- Let say β value is 0.9 then value multiplying derivative will be 0.1.
- However, still there is the smoothening is not achieved as in AdaGrad we have v_{dw} instead of derivative.

Optimizers (RMSProp)

- Implementation Details: On iteration t in minibatch

Compute $\frac{\partial L}{\partial w}$, $\frac{\partial L}{\partial b}$ on current minibatch

$$S_{dw} = \beta S_{dw} + (1-\beta) \left(\frac{\partial L}{\partial w} \right)^2$$

$$S_{db} = \beta S_{db} + (1-\beta) \left(\frac{\partial L}{\partial b} \right)^2$$

$$w_t = w_{t-1} - \eta' \left[\frac{\partial L}{\partial w_{t-1}} \right]$$

$$b_t = b_{t-1} - \eta' \left[\frac{\partial L}{\partial b_{t-1}} \right]$$

Optimizers (Adam)

- Imagine you are walking down a mountain:
 - Momentum → keeps track of your past direction, so you don't zig-zag too much.
 - Adaptive learning rate (RMSProp) → if the slope is steep, take small steps; if it's flat, take bigger steps.
- Adam combines both → smooth + adaptive.

Optimizers (Adam)

- If we need smoothening of gradient as **AdaGrad** and controlled learning rate like **RMSProp** then we can **combine both** to achieve this and such optimizer is called Adam optimizer.
- **Adaptive Moment Estimation (Adam)**
- Here, we used both Sdw and Vdw

$$w_t = w_{t-1} - \eta'_t \frac{\partial L}{\partial w(t-1)}$$

becomes

$$w_t = w_{t-1} - \eta' Vdw$$

$$\eta' = \frac{\eta}{\sqrt{Sdw + \epsilon}}$$

$$vdw_t = \beta \times vdw_{t-1} + (1 - \beta) \times \frac{\partial L}{\partial w}$$

$$Sdw_t = \beta \times Sdw_{t-1} + (1 - \beta) \times \left(\frac{\partial L}{\partial w}\right)^2$$

- Adaptive Learning rate + smoothening
- One of the best optimizer of today with various variants!

Optimizers (Adam)

Compute $\frac{\partial h}{\partial w}$, $\frac{\partial h}{\partial b}$ using current mini batch

$$v_{dw} = \beta_1 v_{dw} + (1 - \beta) \frac{\partial h}{\partial w} \quad \left. \right\} \text{Momentum}$$

$$v_{db} = \beta_1 v_{db} + (1 - \beta) \frac{\partial h}{\partial b} \quad \left. \right\} \text{Momentum}$$

$$s_{dw} = \beta_2 s_{dw} + (1 - \beta) \left(\frac{\partial h}{\partial w} \right)^2 \quad \left. \right\} \text{Rms Prop}$$

$$s_{db} = \beta_2 s_{db} + (1 - \beta) \left(\frac{\partial h}{\partial b} \right)^2 \quad \left. \right\} \text{Rms Prop}$$

$$\boxed{w_t = w_{t-1} - \frac{\eta * v_{dw}}{\sqrt{s_{dw} + \epsilon}} \quad \left. \right\} \text{initial LR}}$$
$$b_t = b_{t-1} - \frac{\eta * v_{db}}{\sqrt{s_{db} + \epsilon}} \quad \left. \right\} \text{Adam optimizer}$$

Summary

- Discussed about loss calculations and its variants
- Discussed about various optimizers
- Discussed about hierarchy of optimizers